

//Assignment 3

// Executing wordmedian mapreduce program

```
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordmedian /hadoopdata/hdfs/test1.txt /hadoopdata/out // command to execute wormedian code.
```

```
18/06/07 22:53:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/06/07 22:53:41 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/06/07 22:53:43 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
18/06/07 22:53:43 INFO input.FileInputFormat: Total input paths to process : 1
18/06/07 22:53:43 INFO mapreduce.JobSubmitter: number of splits:1
18/06/07 22:53:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1528385216097_0003
18/06/07 22:53:44 INFO impl.YarnClientImpl: Submitted application application_1528385216097_0003
18/06/07 22:53:44 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1528385216097_0003/
18/06/07 22:53:44 INFO mapreduce.Job: Running job: job_1528385216097_0003
18/06/07 22:53:55 INFO mapreduce.Job: Job job_1528385216097_0003 running in uber mode : false
18/06/07 22:53:55 INFO mapreduce.Job:  map 0% reduce 0%
18/06/07 22:54:03 INFO mapreduce.Job:  map 100% reduce 0%
18/06/07 22:54:18 INFO mapreduce.Job:  map 100% reduce 100%
18/06/07 22:54:19 INFO mapreduce.Job: Job job_1528385216097_0003 completed successfully
18/06/07 22:54:19 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=86
        FILE: Number of bytes written=215259
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=208
        HDFS: Number of bytes written=32
        HDFS: Number of read operations=6
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=1
        Launched reduce tasks=1
        Data-local map tasks=1
        Total time spent by all maps in occupied slots (ms)=5509
```

```
Total time spent by all reduces in occupied slots (ms)=11188
Total time spent by all map tasks (ms)=5509
Total time spent by all reduce tasks (ms)=11188
Total vcore-milliseconds taken by all map tasks=5509
Total vcore-milliseconds taken by all reduce tasks=11188
Total megabyte-milliseconds taken by all map tasks=5641216
Total megabyte-milliseconds taken by all reduce
tasks=11456512
```

#### Map-Reduce Framework

```
Map input records=4
Map output records=19
Map output bytes=152
Map output materialized bytes=86
Input split bytes=112
Combine input records=19
Combine output records=8
Reduce input groups=8
Reduce shuffle bytes=86
Reduce input records=8
Reduce output records=8
Spilled Records=16
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=156
CPU time spent (ms)=1630
Physical memory (bytes) snapshot=293113856
Virtual memory (bytes) snapshot=4118224896
Total committed heap usage (bytes)=170004480
```

#### Shuffle Errors

```
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
```

#### File Input Format Counters

```
Bytes Read=96
```

#### File Output Format Counters

```
Bytes Written=32
```

The median is: 4

You have new mail in /var/spool/mail/acadgild

// output for wordmedian mapreduce program

```
[acadgild@localhost mapreduce]$ hdfs dfs -ls /hadoopdata/out
18/06/07 22:57:13 WARN util.NativeCodeLoader: Unable to load native-
hadoop library for your platform... using builtin-java classes where
applicable
```

Found 2 items

```
-rw-r--r--  1 acadgild supergroup          0 2018-06-07 22:54
/hadoopdata/out/_SUCCESS
-rw-r--r--  1 acadgild supergroup        32 2018-06-07 22:54
/hadoopdata/out/part-r-00000
```

```
[acadgild@localhost mapreduce]$ hdfs dfs -cat /hadoopdata/out/part-r-00000
```

```
18/06/07 22:57:48 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
1      1
2      6
3      1
4      4
5      1
6      4
7      1
8      1
```

```
You have new mail in /var/spool/mail/acadgild
```

```
// Executing wordmean mapreduce program
```

```
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-
```

```
2.6.5.jar wordmean /hadoopdata/hdfs/test1.txt /hadoopdata/mean-out
```

```
18/06/07 23:00:14 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
18/06/07 23:00:15 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
```

```
18/06/07 23:00:16 INFO input.FileInputFormat: Total input paths to process : 1
```

```
18/06/07 23:00:17 INFO mapreduce.JobSubmitter: number of splits:1
```

```
18/06/07 23:00:17 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1528385216097_0004
```

```
18/06/07 23:00:17 INFO impl.YarnClientImpl: Submitted application application_1528385216097_0004
```

```
18/06/07 23:00:17 INFO mapreduce.Job: The url to track the job:
```

```
http://localhost:8088/proxy/application_1528385216097_0004/
```

```
18/06/07 23:00:17 INFO mapreduce.Job: Running job: job_1528385216097_0004
```

```
18/06/07 23:00:27 INFO mapreduce.Job: Job job_1528385216097_0004 running in uber mode : false
```

```
18/06/07 23:00:27 INFO mapreduce.Job: map 0% reduce 0%
```

```
18/06/07 23:00:34 INFO mapreduce.Job: map 100% reduce 0%
```

```
18/06/07 23:00:55 INFO mapreduce.Job: map 100% reduce 100%
```

```
18/06/07 23:00:57 INFO mapreduce.Job: Job job_1528385216097_0004
```

```
completed successfully
```

```
18/06/07 23:00:57 INFO mapreduce.Job: Counters: 49
```

```
File System Counters
```

```
FILE: Number of bytes read=39
```

```
FILE: Number of bytes written=215439
```

```
FILE: Number of read operations=0
```

```
FILE: Number of large read operations=0
```

```
FILE: Number of write operations=0
```

```
HDFS: Number of bytes read=208
```

```
HDFS: Number of bytes written=19
```

```
HDFS: Number of read operations=6
```

```

        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=1
        Launched reduce tasks=1
        Data-local map tasks=1
        Total time spent by all maps in occupied slots (ms)=4616
        Total time spent by all reduces in occupied slots (ms)=17985
        Total time spent by all map tasks (ms)=4616
        Total time spent by all reduce tasks (ms)=17985
        Total vcore-milliseconds taken by all map tasks=4616
        Total vcore-milliseconds taken by all reduce tasks=17985
        Total megabyte-milliseconds taken by all map tasks=4726784
        Total megabyte-milliseconds taken by all reduce
tasks=18416640
    Map-Reduce Framework
        Map input records=4
        Map output records=38
        Map output bytes=551
        Map output materialized bytes=39
        Input split bytes=112
        Combine input records=38
        Combine output records=2
        Reduce input groups=2
        Reduce shuffle bytes=39
        Reduce input records=2
        Reduce output records=2
        Spilled Records=4
        Shuffled Maps =1
        Failed Shuffles=0
        Merged Map outputs=1
        GC time elapsed (ms)=141
        CPU time spent (ms)=1600
        Physical memory (bytes) snapshot=294711296
        Virtual memory (bytes) snapshot=4117938176
        Total committed heap usage (bytes)=170004480
    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
    File Input Format Counters
        Bytes Read=96
    File Output Format Counters
        Bytes Written=19
The mean is: 4.0
You have new mail in /var/spool/mail/acadgild

// output for wordmean mapreduce program

[acadgild@localhost mapreduce]$ hdfs dfs -ls /hadoopdata/mean-out

```

```
18/06/07 23:01:26 WARN util.NativeCodeLoader: Unable to load native-
hadoop library for your platform... using builtin-java classes where
applicable
Found 2 items
-rw-r--r--    1 acadgild supergroup          0 2018-06-07 23:00
/hadoopdata/mean-out/_SUCCESS
-rw-r--r--    1 acadgild supergroup        19 2018-06-07 23:00
/hadoopdata/mean-out/part-r-00000
[acadgild@localhost mapreduce]$ hdfs dfs -cat /hadoopdata/mean-out/part-
r-00000
18/06/07 23:02:13 WARN util.NativeCodeLoader: Unable to load native-
hadoop library for your platform... using builtin-java classes where
applicable
count 19
length      76
You have new mail in /var/spool/mail/acadgild
```

```
// Executing wordstandarddeviation  mapreduce program
```

```
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-
2.6.5.jar wordstandarddeviation  /hadoopdata/hdfs/test1.txt
/hadoopdata/deviation-out
18/06/07 23:03:40 WARN util.NativeCodeLoader: Unable to load native-
hadoop library for your platform... using builtin-java classes where
applicable
18/06/07 23:03:41 INFO client.RMPProxy: Connecting to ResourceManager at
localhost/127.0.0.1:8032
18/06/07 23:03:43 INFO input.FileInputFormat: Total input paths to
process : 1
18/06/07 23:03:43 INFO mapreduce.JobSubmitter: number of splits:1
18/06/07 23:03:43 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1528385216097_0005
18/06/07 23:03:44 INFO impl.YarnClientImpl: Submitted application
application_1528385216097_0005
18/06/07 23:03:44 INFO mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1528385216097_0005/
18/06/07 23:03:44 INFO mapreduce.Job: Running job: job_1528385216097_0005
18/06/07 23:03:53 INFO mapreduce.Job: Job job_1528385216097_0005 running
in uber mode : false
18/06/07 23:03:53 INFO mapreduce.Job:  map 0% reduce 0%
18/06/07 23:04:00 INFO mapreduce.Job:  map 100% reduce 0%
18/06/07 23:04:11 INFO mapreduce.Job:  map 100% reduce 100%
18/06/07 23:04:15 INFO mapreduce.Job: Job job_1528385216097_0005
completed successfully
18/06/07 23:04:15 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=56
        FILE: Number of bytes written=215643
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=208
```

```

HDFS: Number of bytes written=30
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=4722
  Total time spent by all reduces in occupied slots (ms)=6622
  Total time spent by all map tasks (ms)=4722
  Total time spent by all reduce tasks (ms)=6622
  Total vcore-milliseconds taken by all map tasks=4722
  Total vcore-milliseconds taken by all reduce tasks=6622
  Total megabyte-milliseconds taken by all map tasks=4835328
  Total megabyte-milliseconds taken by all reduce tasks=6780928
Map-Reduce Framework
  Map input records=4
  Map output records=57
  Map output bytes=836
  Map output materialized bytes=56
  Input split bytes=112
  Combine input records=57
  Combine output records=3
  Reduce input groups=3
  Reduce shuffle bytes=56
  Reduce input records=3
  Reduce output records=3
  Spilled Records=6
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=132
  CPU time spent (ms)=1570
  Physical memory (bytes) snapshot=287899648
  Virtual memory (bytes) snapshot=4118224896
  Total committed heap usage (bytes)=170004480
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=96
File Output Format Counters
  Bytes Written=30
The standard deviation is: 2.0
You have new mail in /var/spool/mail/acadgild

```

```
// output for wordstandarddeviation program
```

```
[acadgild@localhost ~]$ hdfs dfs -cat /hadoopdata/deviation-out/part-r-00000
18/06/07 23:15:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
count 19
length      76
square      380
[acadgild@localhost ~]$

[acadgild@localhost mapreduce]$
```